

# ROB-GY 6323

## reinforcement learning and optimal control for robotics

### Lecture I

### Introduction

# Ludovic Righetti

Associate Professor

Mechanical and Aerospace Engineering

Electrical and Computer Engineering

## Contact

[ludovic.righetti@nyu.edu](mailto:ludovic.righetti@nyu.edu)

370 Jay Street, Room 801

## Office Hours

Wednesday 3 to 4pm

Any other time => questions online or by appointment

# Understanding the **fundamental algorithmic principles** of autonomous robotic locomotion and manipulation



Huaijiang  
Zhu



Armand  
Jordana



Avadesh  
Meduri



Sarmad  
Mehrdad



Quim  
Ortiz




Joseph  
Amigo



Quang  
Nam  
Nguyen




Jiangnan  
Zhang




Stacy  
Ashlyn


















Initial policy  
(Zero forces)



**Goals of the class** how to get a robot to move “optimally”

**1 Algorithms** to compute complex robot movements

Using **optimal control** and **reinforcement learning**

**2 Practical application** of algorithms in **real world applications**  
(used in many applications beyond robotics)

## **Necessary background**

- Calculus (derivatives, gradients, Taylor series, etc)
- Linear algebra (positive definiteness, inverse, etc)
- Python



**What is optimal control?**

**What is reinforcement learning?**




## Model-Predictive Control

compute optimal  
trajectories in real-time




## Reinforcement Learning

pre-compute  
a policy offline



[Boston Dynamics 2022]




[Huang et al. 2022]



## Model-Predictive Control


compute optimal  
trajectories in real-time



$$x_{n+1} = x_n + \Delta t \cdot f(x_n, u_n)$$


$$g(x_n, u_n) \leq 0$$

explicit dynamics



## Reinforcement Learning

pre-compute  
a policy offline



$$u = \pi(x, \text{obstacles})$$

$$\min \sum_{n=0}^N l_n(x_n, u_n)$$

**Both approaches solve  
optimal control problems  
=> numerical optimization**



simulator



# Tentative list of topics

- Basics of optimization
- Trajectory optimization:
  - Linear problems (LQR)
  - Nonlinear problems (SQP)
  - Model predictive control
  - Zero order methods (MPPI)
- Policy optimization:
  - Bellman's principle of optimality and dynamic programming
  - Value- and Policy-iteration
  - Deep Q-learning
  - Policy gradients and actor-critic algorithms (PPO)
  - Monte-Carlo tree search



**Your turn...**

**...what do you expect from the class?**



## **Course Project(s) (50%)**

Goal: implement in simulation algorithms seen in class

## **Homework (25%)**

To be handed in every few weeks

(theory + programming exercises in Python)

## **Paper Report (25%)**

## Course material

All necessary material will be posted on Brightspace  
Code will be posted on the Github site of the class

<https://github.com/righetti/optlearningcontrol>

## Discussions/Forum with Slack or Discord?

### Contact

[ludovic.righetti@nyu.edu](mailto:ludovic.righetti@nyu.edu)

Office hours in person

Wednesday 3pm to 4pm

370 Jay street - room 801

(except next week)

### Course Assistant

Armand Jordana

[aj2988@nyu.edu](mailto:aj2988@nyu.edu)

Office hours Monday 1pm to 2pm

Rogers Hall 515



any other time by appointment only



The class will not follow any particular book  
But these resources are good references

## Numerical Optimal Control (Draft)

Sébastien Gros and Moritz Diehl

April 27, 2022

<https://www.syscop.de/files/2020ss/NOC/book-NOCSE.pdf>

## Reinforcement Learning



An Introduction  
second edition

Richard S. Sutton and Andrew G. Barto

Copyrighted Material

<http://incompleteideas.net/book/the-book-2nd.html>

The class will not follow any particular book  
But these resources are good references




D. Bertsekas (Vol. 1+2)



**Questions?**

# Optimal Control Problems date back to XVIIth century

Johann Bernoulli posed the question in 1696: “what is the shape of the ground that allows a ball to reach B from A in the fastest time?”



The optimal curve (red) is called the Brachistochrone curve

Actually, it is not quite yet an optimal control problem

It was originally solved by 5 mathematicians independently

Jacob Bernoulli, Isaac Newton, Gottfried Leibniz, Ehrenfried W. von Tschirnhaus and Guillaume de l'Hôpital





# The theory of optimal control dates from the 1950s



In 1951 H.W. Kuhn and A.W. Tucker formulated first order necessary conditions to find a minimum of a nonlinear optimization problem - W. Karush had found them in his M.Sc. thesis in 1939



In 1954, Richard Bellman formulated a “principle of optimality” to provide necessary conditions to solve optimal control problems (aka the dynamic programming algorithm)



In 1956, Lev Semyonovich Pontryagin formulated a “maximum principle” to provide necessary conditions to solve optimal control problems

These ideas are the foundations of most modern optimal control and reinforcement learning algorithms

# **Optimal control $\simeq$ Reinforcement learning**

Finding the “best” way to solve a  
sequential decision making problem




A sequential decision making problem





# A sequential decision making problem





# A sequential decision making problem

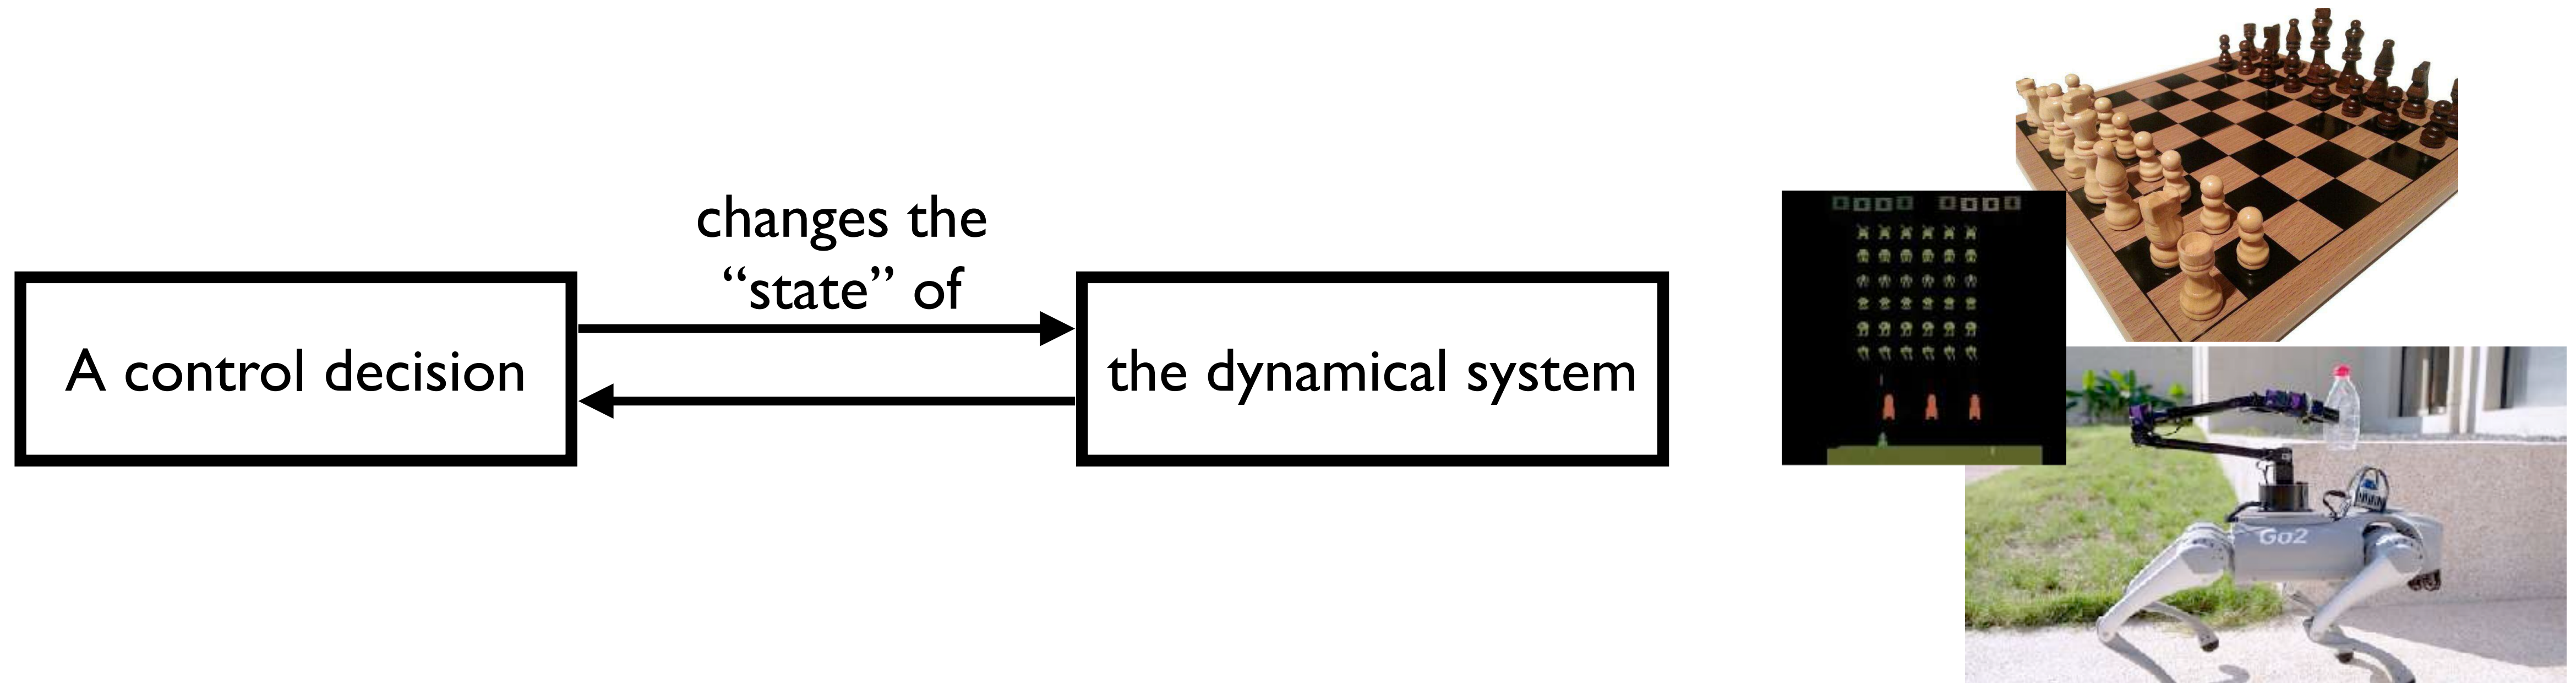


Traversing a series of challenging terrains

Luo et al. 2024




# A sequential decision making problem



The problem: find the “**best sequence of actions**” to make the **dynamical system** behave as desired (e.g. win the game or do a backflip)




# Best sequence of actions?



**Goal**  
**X**

- How to accelerate a car to reach a goal in minimum time?
- How to accelerate a car to reach a goal in minimum acceleration?
- How to accelerate a car to minimize fuel consumption?
- How to accelerate a car to maximize passenger comfort?


minimum time?



**Goal**  
**X**





# minimum acceleration?




**Goal**  
**X**

# Different measures of “best” lead to very different answers

minimum time?




minimum acceleration?



**Goal**  
**X**




# A sequential decision making problem



The problem: find the “**best sequence of actions**” to make the **dynamical system** behave as desired (e.g. win the game or do a backflip)



# Dynamical systems are all over the place



The state of a dynamical system changes according to some “rules”

- The rules of the game and the opponent moves (chess)
- Newton’s law of motion (robot)
- Energy production rate and consumption (smart grid)
- etc



# Robots are continuous time dynamical systems



# Discrete vs. continuous time dynamical systems





# Discrete vs. continuous time dynamical systems

# A first optimal control problem





# Structure of an optimal control problem

$$\min_{u_0, u_1, \dots, u_N} \sum_{i=0}^N g_i(x_i, u_i)$$

Find actions that  
optimize a  
performance cost

Subject to:

$$x_{n+1} = f(x_n, u_n)$$

$$h_n(x_n, u_n) \leq 0$$

to control a  
dynamical system  
(maybe with  
constraints)

This is an optimization problem

# **Basics of optimization**



# Off-the-shelf optimization algorithms



Fundamental algorithms for scientific computing in Python

GET STARTED

```
scipy.optimize.minimize(fun, x0, args=(), method=None, jac=None, hess=None,
                        hessp=None, bounds=None, constraints=(), tol=None, callback=None, options=None)
```


## Optimization (`scipy.optimize`)

- Unconstrained minimization of multivariate scalar functions (`minimize`)
  - Nelder-Mead Simplex algorithm (`method='Nelder-Mead'`)
  - Broyden-Fletcher-Goldfarb-Shanno algorithm (`method='BFGS'`)
    - Avoiding Redundant Calculation
  - Newton-Conjugate-Gradient algorithm (`method='Newton-CG'`)
    - Full Hessian example:
    - Hessian product example:
  - Trust-Region Newton-Conjugate-Gradient Algorithm (`method='trust-ncg'`)
    - Full Hessian example:
    - Hessian product example:
  - Trust-Region Truncated Generalized Lanczos / Conjugate Gradient Algorithm (`method='trust-krylov'`)
    - Full Hessian example:
    - Hessian product example:
  - Trust-Region Nearly Exact Algorithm (`method='trust-exact'`)
- Constrained minimization of multivariate scalar functions (`minimize`)
  - Trust-Region Constrained Algorithm (`method='trust-constr'`)
    - Defining Bounds Constraints:
    - Defining Linear Constraints:
    - Defining Nonlinear Constraints:
    - Solving the Optimization Problem:
  - Sequential Least Squares Programming (SLSQP) Algorithm (`method='SLSQP'`)
- Global optimization
- Least-squares minimization (`least_squares`)
  - Example of solving a fitting problem
  - Further examples
- Univariate function minimizers (`minimize_scalar`)
  - Unconstrained minimization (`method='brent'`)
  - Bounded minimization (`method='bounded'`)
- Custom minimizers
- Root finding
  - Scalar functions
  - Fixed-point solving
  - Sets of equations
  - Root finding for large problems
  - Still too slow? Preconditioning.
- Linear programming (`linprog`)
  - Linear programming example
- Assignment problems
  - Linear sum assignment problem example
- Mixed integer linear programming
  - Knapsack problem example

# Minimizing a function

$$\min_x x^2$$

$$x^* = 0$$






# Minimizing a function

$$\min_x 5x^2 - 10x$$

$$x^* = 1$$



# Minimums (local and global)




These conditions are not sufficient!

# Optimization with constraints

$$\min_x 5x^2 - 10x$$

$$x < 0$$

$$x^* = 0$$






# Optimization with constraints

$$\min_x 5x^2 - 10x$$

$$2x + 4 = 0$$

$$x^* = -2$$




# Optimization with constraints

$$\min_x 5x^2 - 10x$$

$$2x + 4 = 0$$


$$-x < 0$$

Unfeasible



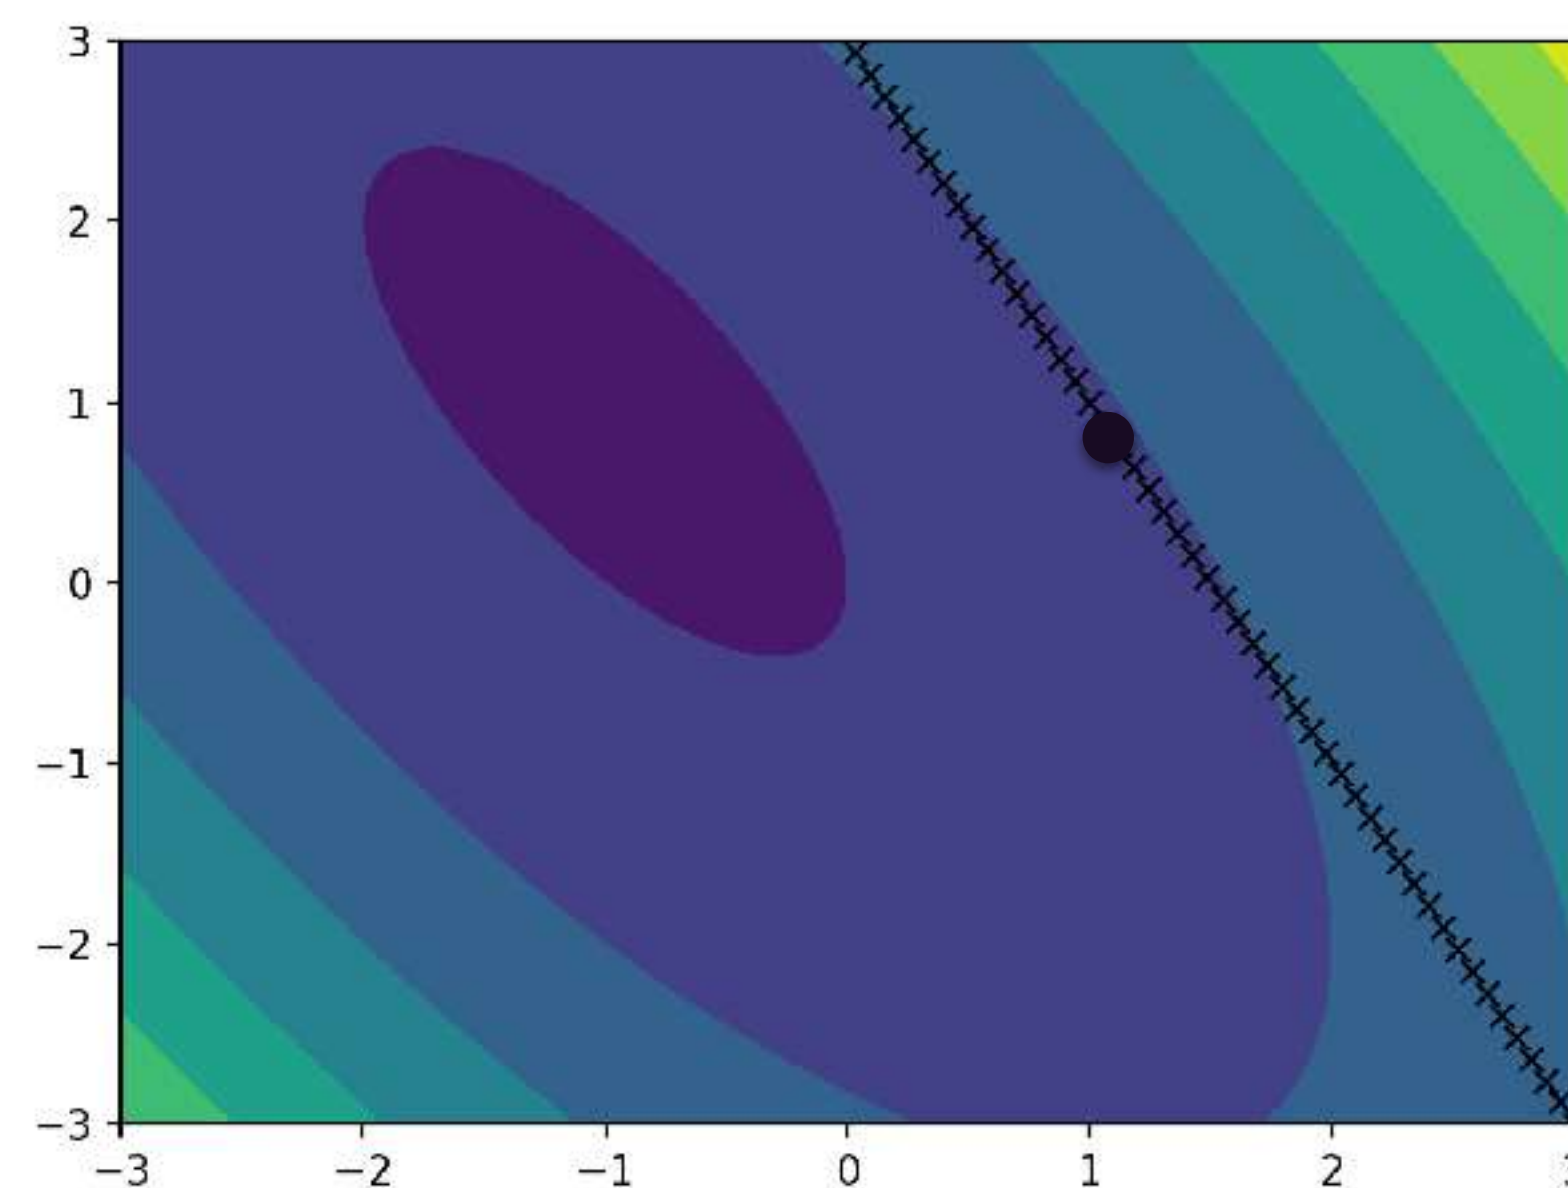


$$\min_{x_1, x_2} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}^T \begin{bmatrix} 2 & 1 \\ 1 & 1 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} + 2x_1$$



$$\min_{x_1, x_2} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}^T \begin{bmatrix} 2 & 1 \\ 1 & 1 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} + 2x_1$$

**s.t.**  $x_1 + 2x_2 + 3 = 0$



$$\min_{x_1, x_2} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}^T \begin{bmatrix} 2 & 1 \\ 1 & 1 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} + 2x_1$$

**s.t.**  $x_1 + 2x_2 + 3 = 0$   
 $x_1 < 0.5$

