

ROB-GY 6323

reinforcement learning and optimal control for robotics

Lecture 9

Q-learning

Course material

All necessary material will be posted on Brightspace
Code will be posted on the Github site of the class

<https://github.com/righetti/optlearningcontrol>

Discussions/Forum with Slack

Contact

ludovic.righetti@nyu.edu

Office hours in person
Wednesday 3pm to 4pm
370 Jay street - room 801

Course Assistant

Armand Jordana
aj2988@nyu.edu

Office hours Monday 1pm to 2pm
Rogers Hall 515



any other time by appointment only

Tentative schedule (subject to change)

Week	Lecture		Homework	Project
1	<u>Intro</u>	Lecture 1: introduction		
2	<u>Trajectory optimization</u>	Lecture 2: Basics of optimization	HW 1	
3		Lecture 3: QPs		
4		Lecture 4: Nonlinear optimal control		
5		Lecture 5: Model-predictive control		
6		Lecture 6: Sampling-based optimal control	HW 2	
7	Lecture 7: Bellman's principle			
8	<u>Policy optimization</u>	Lecture 8: Value iteration / policy iteration		Project 1
9		Lecture 9: Q-learning	HW 3	
10		Lecture 10: Deep Q learning		
11		Lecture 11: Actor-critic algorithms		
12		Lecture 12: Learning by demonstration	HW 4	Project 2
13	Lecture 13: Monte-Carlo Tree Search			
14		Lecture 14: Beyond the class		
15	Finals week			

Infinite horizon problems

In general the sum of costs might diverge

We introduce discounted costs $\lim_{N \rightarrow \infty} \min_{\pi(x_n)} \sum_{n=0}^N \alpha^n l(x_n, \pi(x_n))$

$0 < \alpha < 1$ Discount factor

This will work as long as $l(x_n, \pi(x_n))$ is bounded: $|l(x, u)| < M$

Value iteration

For any bounded function $J(x)$, the iteration

$$J^{n+1}(x) = \min_u g(x, u) + \alpha J^n(f(x, u))$$

with $J^0(x) = J(x)$ converges to the optimal value function, i.e.

$$\lim_{n \rightarrow \infty} J^{n+1}(x) = J^*(x)$$

Value iteration algorithm: start from an arbitrary $J(x)$ and iterate $J^{n+1}(x)$

Policy evaluation: how good is a policy?

Policy evaluation algorithm I

For any bounded function $J(x)$, the iteration

$$J_{\mu}^{n+1}(x) = g(x, \mu(x)) + \alpha J_{\mu}^n(f(x, \mu(x)))$$

converges to the total cost of the stationary policy $\mu(x)$, i.e.

$$\lim_{n \rightarrow \infty} J_{\mu}^n(x) = J_{\mu}(x)$$

Policy evaluation algorithm: start from an arbitrary $J(x)$ and iterate $J_{\mu}^{n+1}(x)$

Policy Evaluation Algorithm II

Given a policy μ :

$$\text{Solve } (I - \alpha A)J_\mu = \bar{g}$$

J_μ can also be written as a vector $\begin{pmatrix} J_\mu(x_1) \\ J_\mu(x_2) \\ \vdots \\ J_\mu(x_N) \end{pmatrix}$

Let \bar{g} the vector of costs for all the states, i.e. $\bar{g} = \begin{pmatrix} g(x_1, \mu(x_1)) \\ g(x_2, \mu(x_2)) \\ \vdots \\ g(x_N, \mu(x_N)) \end{pmatrix}$

$J_\mu(f(x, \mu(x)))$ can be written using the matrix A of transitions such that we get the following linear equation

$$J_\mu = \bar{g} + \alpha A J_\mu$$

Policy iteration: finding the optimal policy

Policy Iteration algorithm: optimal policy through policy evaluation

Start with an initial guess for the policy μ_0

1. Policy evaluation step:

Compute $J_{\mu_n}(x)$ using the policy evaluation algorithm

2. Policy update step:

Update the policy using

$$\mu_{k+1} = \arg \min_u g(x, u) + \alpha J_{\mu_k}(f(x, u))$$

Iterate until convergence
(guaranteed to happen in a finite number of iteration!)

Checking all the possible x is not always possible
Can we do something less “perfect” but more practical?

=> reinforcement learning!

A historical digression on reinforcement learning

What is reinforcement learning?

Learning “how to act” from direct interaction with the environment with the goal to “maximize” a reward

=> aka optimal control without a model

=> original RL was supposed to work on real robots


Not a new field!

Early versions of TD-learning from the 1950s

Q-learning in the 1980s

RL with neural networks in the 1990s

$$\max_{u_n} \sum_{n=0}^N R_n(x_n, u_n)$$




$$x_{n+1} = f(x_n, u_n)$$

Markov Decision Process

Early model-based reinforcement learning in robotics

[Schaal and Atkeson ~1995]



Early model-based reinforcement learning in robotics

[Schaal and Atkeson ~1995]

Model-based Reinforcement Learning of Devilsticking

Stefan Schaal & Chris Atkeson



Apprenticeship learning

[Abbeel, 2010]



Apprenticeship learning

[Abbeel, 2010]



What is reinforcement learning?

Today RL algorithms mostly uses simulators (i.e. models) and are data intensive
RL algorithms are really just optimization algorithms to find policies

RL algorithms are really just optimization algorithms to find policies
Deep learning has unlocked the ability to approximate complex policies
=> RL can be applied to very complex problems



[Silver et al., 2016]



[Hwangbo et al. 2019]

Typical RL problems



Set of actions is discrete
State is discrete (countable)

Robotics RL problems



State is continuous
Action space is continuous

Most methods designed for discrete state/action models do not carry over to continuous state/action models

Some notations

In the RL literature

States S_n

Control inputs are called actions A_n

Non-deterministic systems are considered

$$x_{n+1} = f(x_n, u_n, \omega_n)$$

So the optimization criteria is often

$$\min_{u_n} \mathbb{E}_{\omega_n} \left(\sum_{n=0}^{N-1} g_n(x_n, u_n) \right)$$

RL algorithms are often “sampling based”

Running an episode

Idea: try some policy and record the states and instantaneous costs



execute a policy guess on the robot
(or on a simulator) and collect data

update the value/policy guess based on the data

Remember: trajectory optimization can also be sampling based!

Single shooting with sampling based gradient descent


$$\min_{u_0, \dots, u_{N-1}} \sum_{n=0}^{N-1} l_n(f^n(x_0, u_0, \dots, u_{n-1}), u_n) + l_N(f^N(x_0, u_0, \dots, u_{N-1}))$$

Start with a control guess $\bar{u} = [u_0, \dots, u_{N-1}]$, then repeat until convergence

- Sample N trajectories $\bar{u} + \epsilon_n$ where $\epsilon \in \mathcal{N}(0, I)$
- Estimate the gradient of the cost $\nabla l(\bar{u}) \simeq \frac{1}{N\sigma} \sum (f(\bar{u} + \epsilon_n) - f(\bar{u})) \epsilon_n$
- Update $\bar{u} \leftarrow \bar{u} + \alpha \nabla l(\bar{u})$

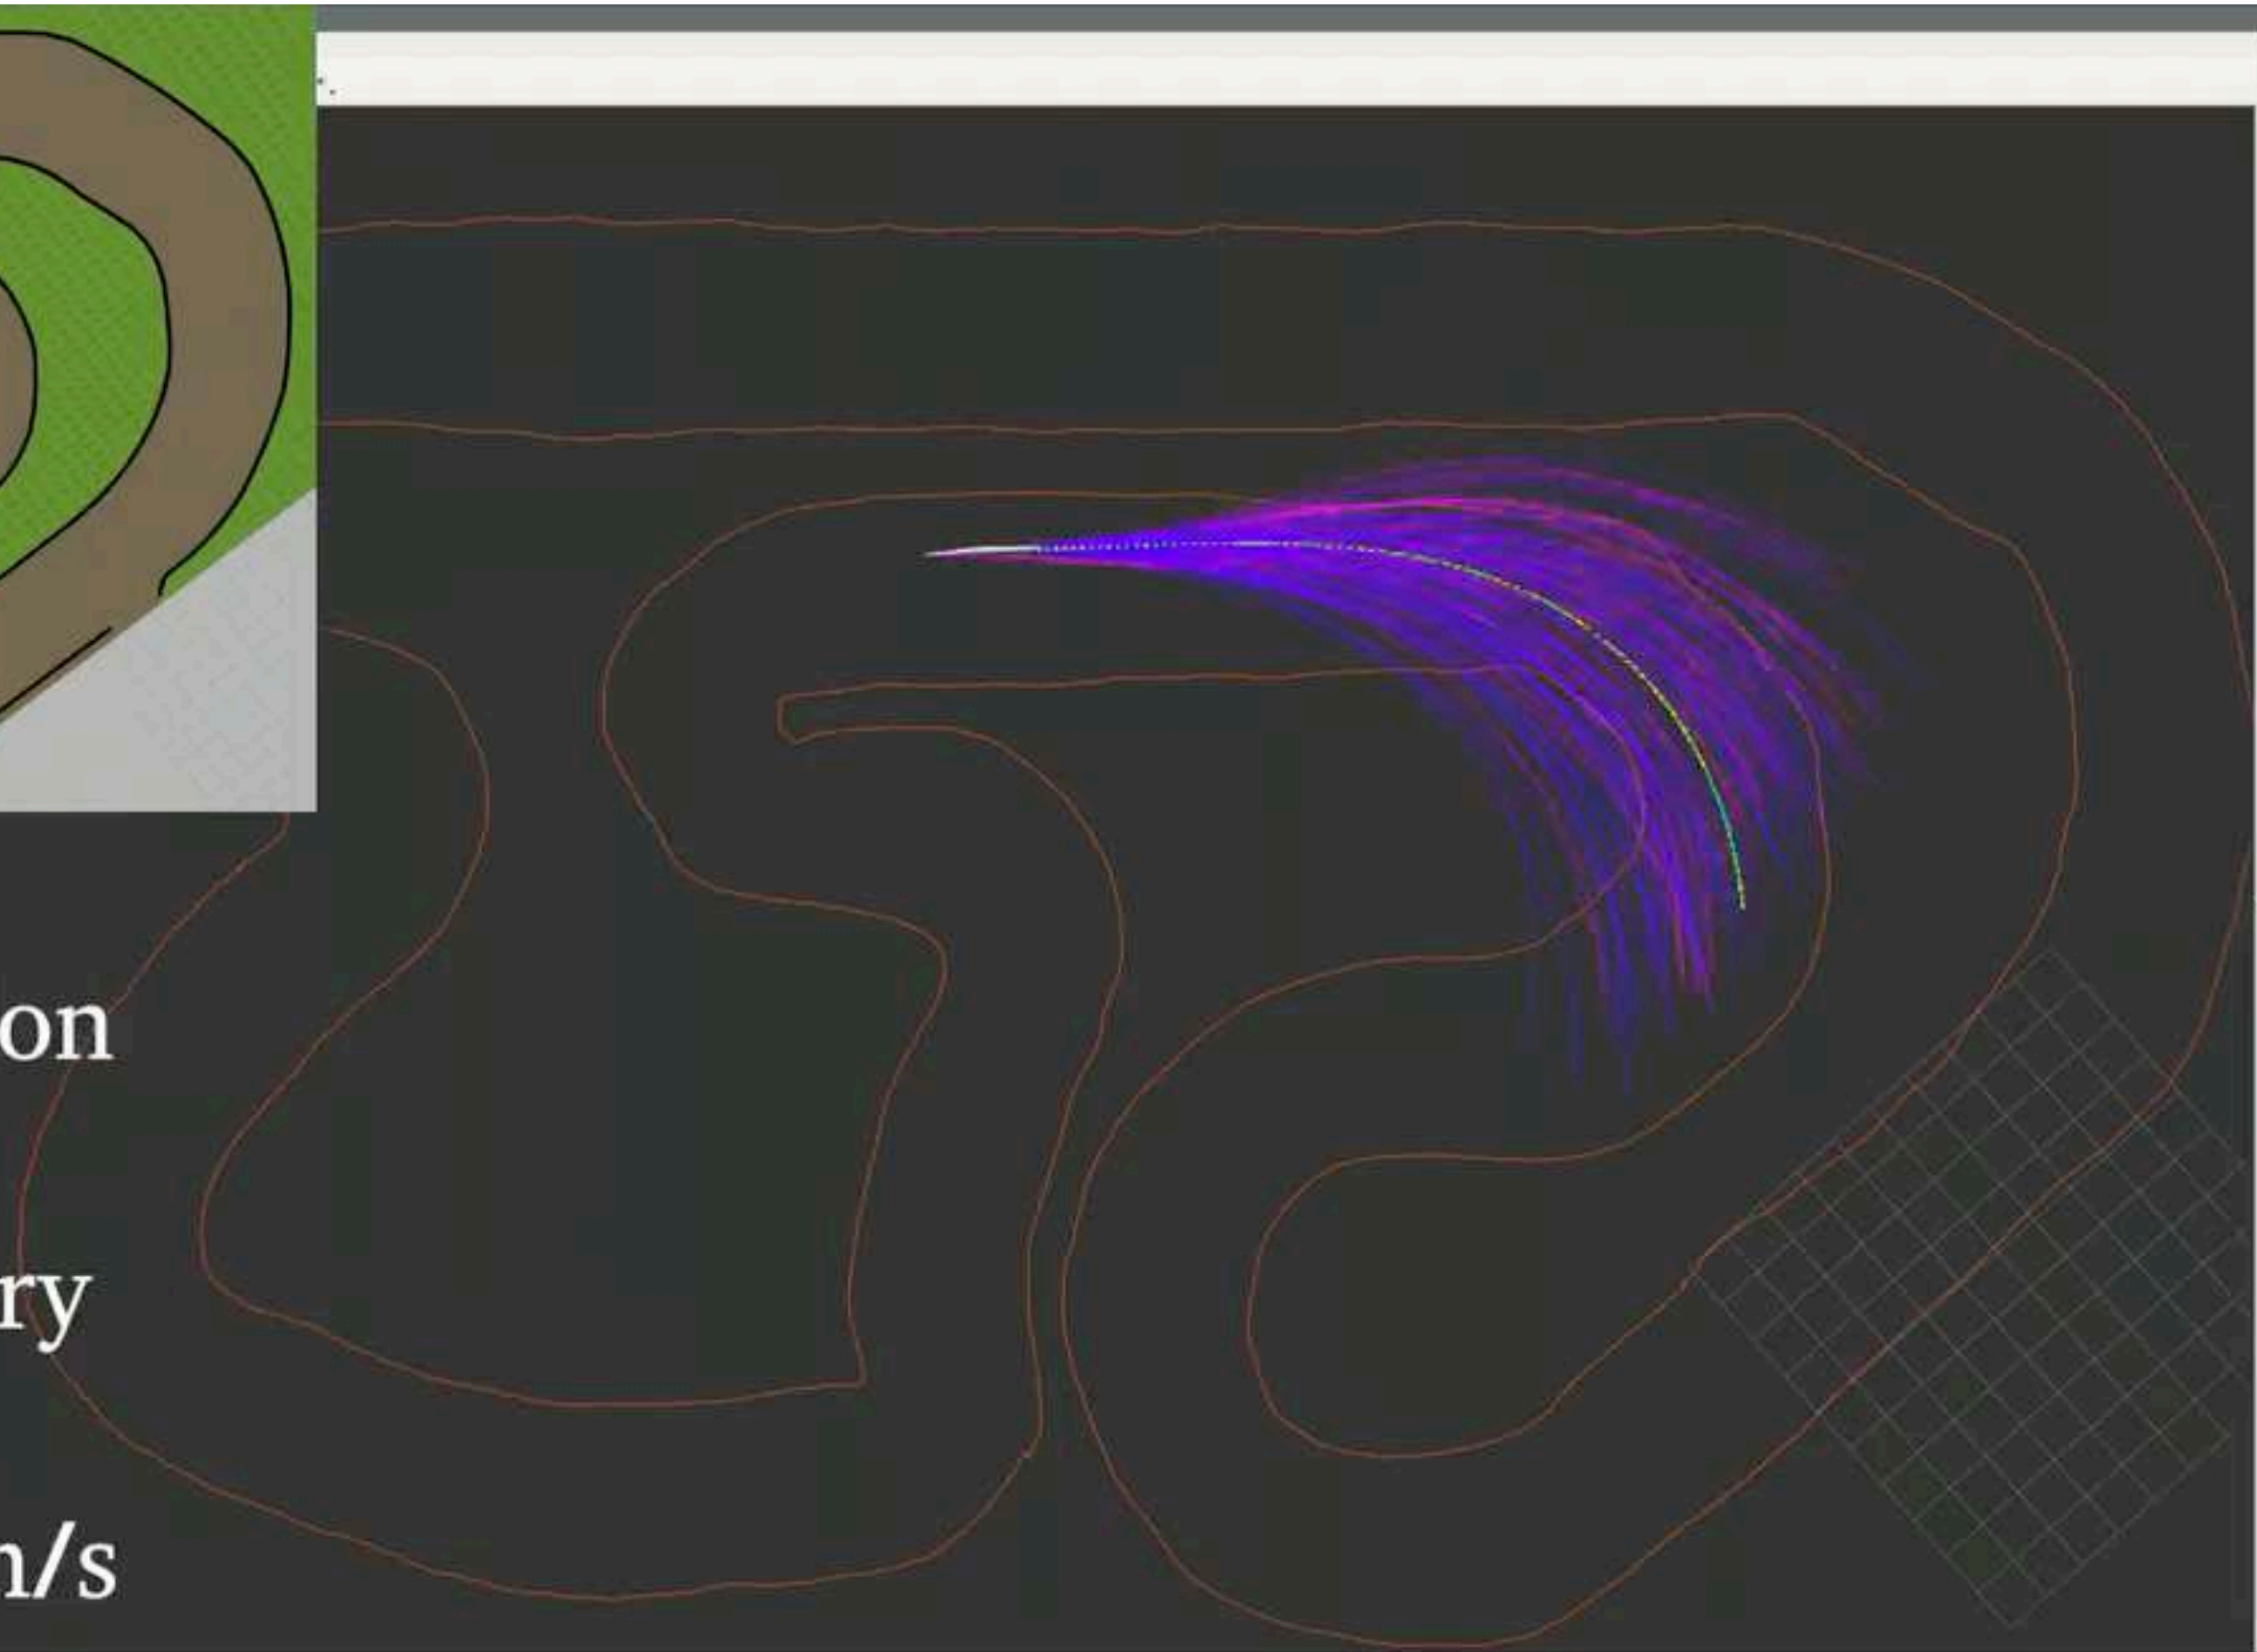
Very convenient as we can replace the dynamics $f()$ by a simulator!

MPPI (Model-predictive path integral control)



Notice the
feedback application
in sampling
results in recovery

Robust MPPI 9 m/s



Experiment 2: Pen Grasping

In order to pick up a pen lying on a table, we learn the downward force profile required to ensure successful grasping of the pen.

[Kalakrishnan et al. 2011]

Trajectory optimization with real-world data

Remember: trajectory optimization can also be sampling based!

Most modern RL algorithms optimize a policy

On-policy vs. off-policy methods

On-policy methods evaluate or improve the policy that is used to generate episodes (i.e. learn while you act)

Off-policy methods evaluate or improve a policy different from the method used to generate the data / episodes.

Before optimizing a policy... we evaluate the value of
a policy

Policy evaluation with sampling I

Monte-Carlo methods

Policy evaluation with Monte-Carlo methods

Policy evaluation with sampling II

TD-learning

Optimal Value function

$$J = \min_u g(x, u) + \alpha J(f(x, u))$$

Value Function for Policy $\mu(x) : x \rightarrow u$

$$J_\mu(x) = g(x, \mu(x)) + \alpha J_\mu(f(x, \mu(x)))$$

Action-value function $Q(x_t, u_t)$

$$Q(x_t, u_t) = g(x_t, u_t) + \alpha \min_u Q(x_{t+1}, u)$$

$$J_{\mu}(x_t) \stackrel{?}{=} \underbrace{g(x_t, \mu(x_t)) + \alpha J_{\mu}(x_{t+1})}_{\text{actual return following time t}}$$

current guess
of cost for state x_t

Temporal difference
error (TD error)

$$\delta_t = g(x_t, \mu(x_t)) + \alpha J_{\mu}(x_{t+1}) - J_{\mu}(x_t)$$

Temporal difference learning

[Sutton 1988]

[Samuel 1959]

TD(0) learning for estimating J_μ

Input: policy to be evaluated μ

Choose a step size $\gamma \in [0, 1]$

Initialize J_μ for all states x

For each episode of length N :

 Choose an initial state x_0


 Loop for each step of the episode:

 Do $\mu(x_t)$

 Observe x_{t+1} Compute $g(x_t, \mu(x_t))$

 Update $J_\mu(x_t) \leftarrow J_\mu(x_t) + \gamma \delta_t$

 using $\delta_t = g(x_t, \mu(x_t)) + \alpha J_\mu(x_{t+1}) - J_\mu(x_t)$



How can we improve the policy?

Q-learning: off-policy TD control

[Watkins 1989]

Action-value function $Q(x_t, u_t)$

$$Q(x_t, u_t) = g(x_t, u_t) + \alpha J^*(x_{t+1})$$

$Q(x_t, u_t)$: cost of doing u_t at state x_t and behaving optimally after

$$J^*(x_t) = \min_u Q(x_{t+1}, u) \quad \text{Optimal value function}$$

$$\mu^*(x) = \arg \min_u Q(x_t, u) \quad \text{Optimal policy}$$

Q-learning: off-policy TD control

[Watkins 1989]

Action-value function $Q(x_t, u_t)$

$$Q(x_t, u_t) = g(x_t, u_t) + \alpha \min_u Q(x_{t+1}, u)$$

$$J^*(x_t) = \min_u Q(x_{t+1}, u) \quad \text{Optimal value function}$$

$$\text{TD-error} \quad \delta_t = g(x_t, u_t) + \alpha \min_u Q(x_{t+1}, u) - Q(x_t, u_t)$$

$$\text{Update} \quad Q(x_t, u_t) \leftarrow Q(x_t, u_t) + \gamma \delta_t$$

Q-learning

$$\delta_t = g(x_t, u_t) + \alpha \min_u Q(x_{t+1}, u) - Q(x_t, u_t)$$

$Q(x, u)$ is stored as a table

	$x = 0$	$x = 0.1$	$x = 0.2$		$x = 6.2$
$u = -5$	$Q(x_1, u_1)$			• • •	
$u = -4.9$				• • •	
				• • •	
	• • •	• • •		• • •	
$u = 0$					
	• • •	• • •		• • •	
$u = 5$				• • •	

The exploration/exploitation trade-off

How do we choose the policy in an episode?

Current optimal guess: $u_t = \arg \min_u Q(x_t, u)$

If we always choose the optimal guess, we might miss better actions/states that we would never try

If we only choose random actions, we might be not be able to get a good guess for Q

ϵ -greedy policy

$$u_t = \begin{cases} \arg \min_u Q(x_t, u) & \text{with probability } 1 - \epsilon \\ \text{random action} & \text{with probability } \epsilon \end{cases}$$

Q-learning

[Watkins 1989]

Choose a step size $\gamma \in [0, 1]$ and small ϵ

Initialize $Q(x, u)$ for all states x and actions u

For each episode:

 Choose an initial state x_0

 Loop for each step of the episode:



 Choose an action using an ϵ -greedy policy from Q

 Observe x_{t+1} Compute $g(x_t, \mu(x_t))$

 Update $Q(x_t, u_t) \leftarrow Q(x_t, u_t) + \gamma \delta_t$

 using $\delta_t = g(x_t, u_t) + \alpha \min_u Q(x_{t+1}, u) - Q(x_t, u_t)$

Q-learning



[source: <https://www.youtube.com/watch?v=YLAWnYAsai8>]

Q-learning

Model-free approach to learn optimal policies
(value iteration with a twist)

$$Q(x_t, u_t) \leftarrow Q(x_t, u_t) + \gamma \delta_t$$

$$\delta_t = g(x_t, u_t) + \alpha \min_u Q(x_{t+1}, u) - Q(x_t, u_t)$$

Guaranteed to converge at infinity BUT can take a long time!

Need to store Q / discrete actions and states


Need to compute the min of Q (expensive!)

Typical RL problems



Set of actions is discrete
State is discrete (countable)

Robotics RL problems



State is continuous
Action space is continuous

Most methods designed for discrete state/action models do not carry over to continuous state/action models

Q-learning





$$\delta_t = g(x_t, u_t) + \alpha \min_u Q(x_{t+1}, u) - Q(x_t, u_t)$$

Can we get rid of the discrete states and actions
and the table?

$Q(x, u)$ is a function - can we approximate it?

	$x = 0$	$x = 0.1$	$x = 0.2$		$x = 6.2$
$u = -5$	$Q(x_1, u_1)$			$\cdot \cdot \cdot$	
$u = -4.9$				$\cdot \cdot \cdot$	
				$\cdot \cdot \cdot$	
	\vdots	\vdots		$\cdot \quad \cdot \quad \cdot$	
$u = 0$					
	\vdots	\vdots		$\cdot \quad \cdot \quad \cdot$	
$u = 5$				$\cdot \cdot \cdot$	

Function approximation



Linear least squares

Given N data point $(x_1, y_1) (x_2, y_2) \cdots (x_N, y_N)$

$y = \sum_{k=0}^K a_k x^k$ Find a function that is a linear combination
of polynomials of the input x

Minimize the least square error between
the output data and the function

$$\min_{a_0 \cdots a_K} \sum_{i=0}^{N-1} \left(\sum_{k=0}^K a_k x_i^k - y_i \right)^2$$

Linear least squares

Minimize the least square error between the output data and the function

$$\min_{a_0 \cdots a_K} \sum_{i=0}^{N-1} \left(\sum_{k=0}^K a_k x_i^k - y_i \right)^2$$

We can write these relations in matrix form by noticing that

$$\sum_{k=0}^K a_k x_i^k = \begin{bmatrix} 1 & x_i & x_i^2 & \cdots & x_i^K \end{bmatrix} \begin{bmatrix} a_0 \\ a_1 \\ a_2 \\ \vdots \\ a_K \end{bmatrix}$$

Linear least squares

We can write these relations in matrix form by noticing that

$$\sum_{k=0}^K a_k x_i^k = \begin{bmatrix} 1 & x_i & x_i^2 & \cdots & x_i^K \end{bmatrix} \begin{bmatrix} a_0 \\ a_1 \\ a_2 \\ \vdots \\ a_K \end{bmatrix}$$

Using all the data points and the knowledge of the degree K and we can then construct the $N \times K$ matrix

$$X = \begin{bmatrix} 1 & x_0 & x_0^2 & \cdots & x_0^K \\ 1 & x_1 & x_1^2 & \cdots & x_1^K \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ 1 & x_{N-1} & x_{N-1}^2 & \cdots & x_{N-1}^K \end{bmatrix}$$

and the vector

$$Y = \begin{bmatrix} y_0 \\ y_1 \\ \vdots \\ y_{N-1} \end{bmatrix} \qquad a = \begin{bmatrix} a_0 \\ a_1 \\ \vdots \\ a_K \end{bmatrix}$$

where each row i of X and Y is defined by the sample i from the dataset.

Linear least squares

We now have

$$Xa - Y = \begin{bmatrix} \sum_{k=0}^K a_k x_0^k - y_0 \\ \sum_{k=0}^K a_k x_1^k - y_1 \\ \vdots \\ \sum_{k=0}^K a_k x_{N-1}^k - y_{N-1} \end{bmatrix}$$

and the original problem can be written as

$$\min_a (Xa - Y)^T (Xa - Y)$$

which is equal to

$$\min_a a^T X^T X a - 2Y^T X a + Y^T Y$$


$$\frac{\partial}{\partial a} (a^T X^T X a - 2Y^T X a + Y^T Y) = 2X^T X a - 2X^T Y = 0$$

$$a = (X^T X)^{-1} X^T Y$$

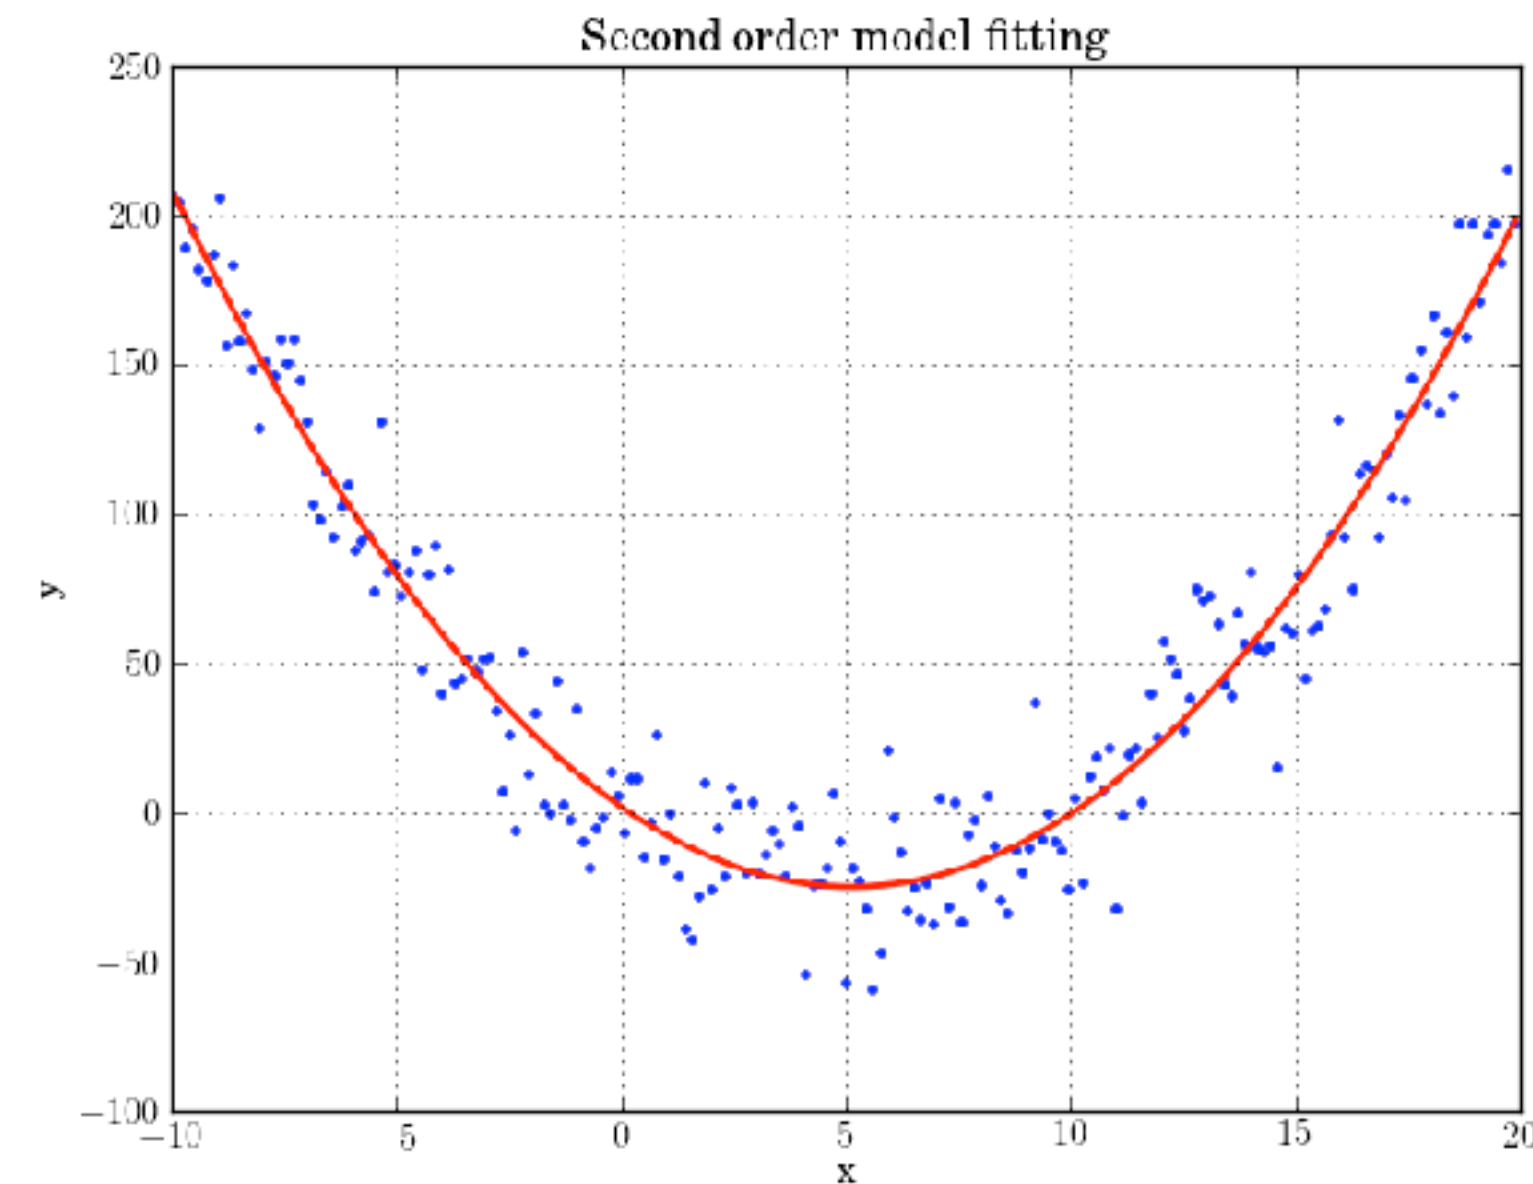
Function approximation

500 noisy samples $\mathcal{D} = \{(x_0, y_0), (x_1, y_1), \dots, (x_N, y_N)\}$

from a quadratic function



$$y = a_1x + a_0$$



$$y = a_2x^2 + a_1x + a_0$$

Nonlinear least-square

$$y = f(x, w) \qquad \min \sum_{i=0}^N (y_i - f(x_i, \omega))^2$$



we cannot compute w explicitly as before
because the function is not linear in w anymore

=> do gradient descent to minimize the function

$$\omega \leftarrow \omega - \eta \frac{\partial}{\partial \omega} \left(\sum_{i=0}^N (y_i - f(x_i, \omega))^2 \right) = \omega + 2\eta \sum_{i=0}^N \left((y_i - f(x_i, \omega)) \frac{\partial f}{\partial \omega} \right)$$

Example: a simple nonlinear function


$$y = \frac{1}{1 + e^{-wx}} \quad \text{a sigmoid function}$$



Example: a simple nonlinear function

multiple inputs and one output

$$y = \frac{1}{1 + e^{-\sum_k w_k x_k}}$$




w_k weights (unknown parameters to find or “learn”)

Example: a simple nonlinear function

multiple inputs and one output

$$y = \frac{1}{1 + e^{-\sum_k w_k x_k}}$$



$$\min_w \sum_{n=0}^N \left(y_i - \frac{1}{1 + e^{-\sum_k w_k x_{k,i}}} \right)$$

Example: a simple nonlinear function



$$\min_w \sum_{n=0}^N \left(y_i - \frac{1}{1 + e^{-\sum_k w_k x_{k,i}}} \right)$$

Dataset has multidimensional input

$$\mathcal{D} = \{(x_{0,0}, x_{1,0}, \dots, x_{M,0}, y_0), (x_{0,1}, x_{1,1}, \dots, x_{M,1}, y_1), \dots, (x_{0,N}, x_{1,N}, \dots, x_{M,N}, y_N)\}$$

Gradient descent enables the optimization of the unknown parameters (or neuron weights) w


One layer neural network



Simple neural network: a combination of neurons to create a (complex) nonlinear function


Deep neural network

A neural network with many layers





Example of a fully connected deep neural network

Multi-dimensional output



Different activation functions



we can replace the sigmoid by other nonlinear functions, popular ones are

- Rectified Linear Units (ReLU) $y = \max\{0, \sum_k w_k x_k\}$
- Softplus $y = \log(1 + e^{\sum_k w_k x_k})$

Different architectures

We can decide:

- How many layers
- How many neurons per layer
- Which activation functions are used
- How layers connect to each other
- etc

Stochastic gradient descent

Usually we cannot do gradient descent using all the data point available in our dataset!

Stochastic gradient descent: randomly select a small number of data points (a mini-batch) and do gradient descent using these points

Libraries to “train” neural networks



Example

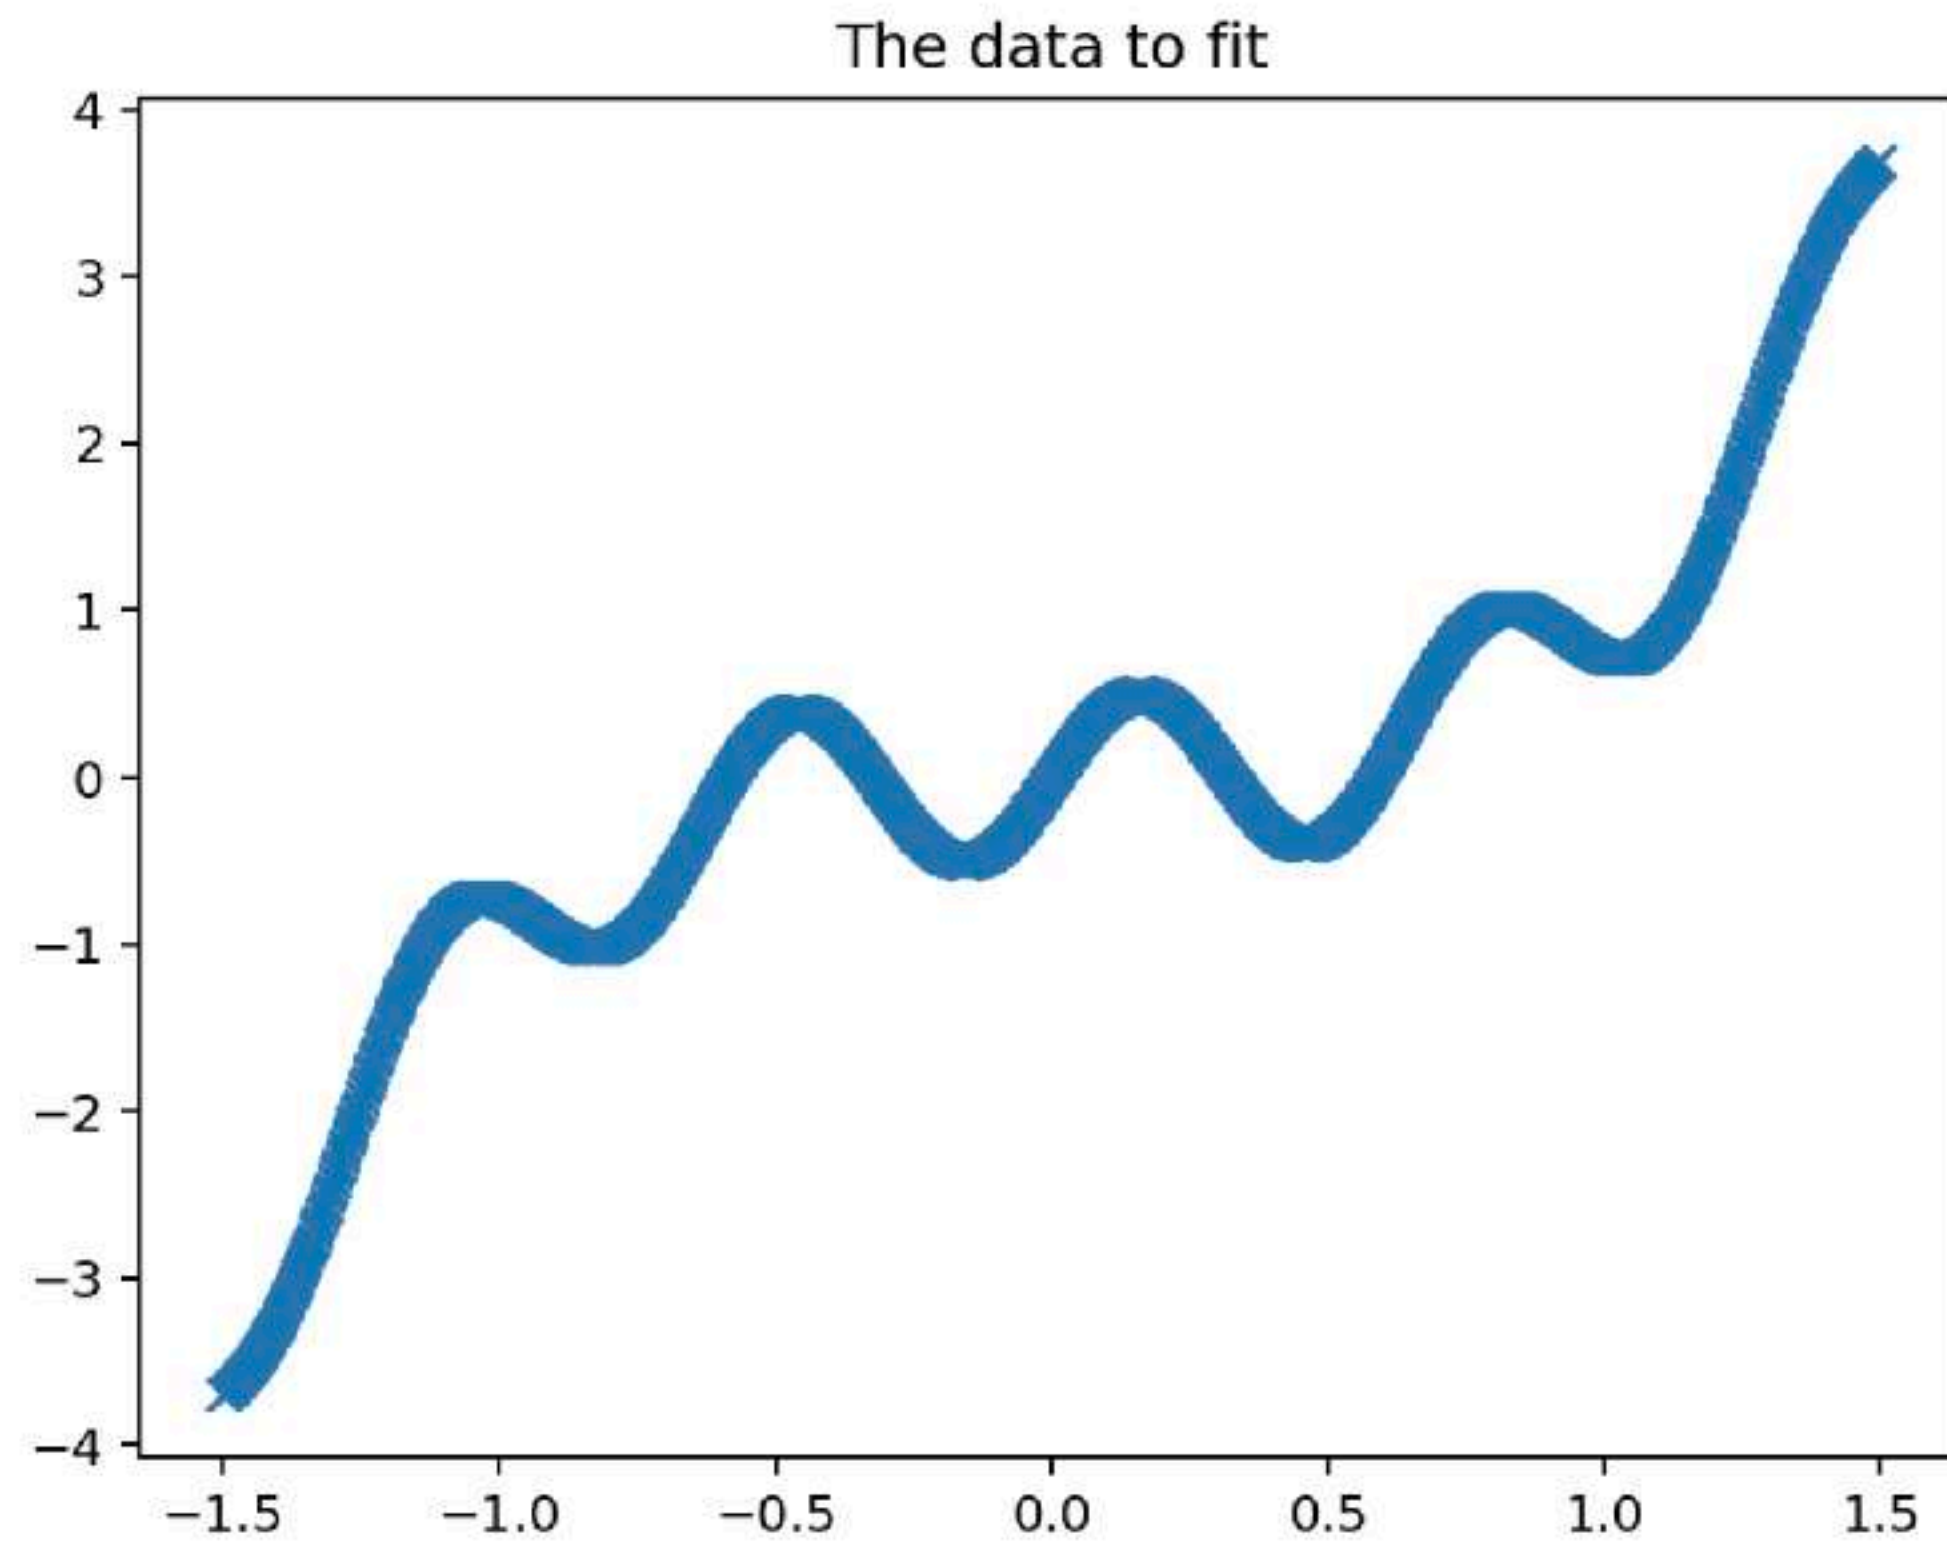
```
In [1]: %matplotlib notebook

import torch

import numpy as np
import matplotlib as mp
import matplotlib.pyplot as plt
```

```
In [2]: ## we create the data to learn from
N = 1000 # number of data points
x = torch.linspace(-1.5, 1.5, steps=N).reshape(N, 1)
y = x**3 + 0.5*torch.sin(10*x)

plt.figure()
plt.plot(x.numpy(), y.numpy(), 'x')
plt.title('The data to fit')
```



In [6]:

```
# # we create another model with 3 hidden layers
model = torch.nn.Sequential(
    torch.nn.Linear(D_in, H),
    torch.nn.ReLU(),
    torch.nn.Linear(H, H),
    torch.nn.ReLU(),
    torch.nn.Linear(H, H),
    torch.nn.ReLU(),
    torch.nn.Linear(H, D_out),
)

# we define the learning rate and select an optimizer
learning_rate = 1e-3
optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)

# we learn doing 5000 iterations
for t in range(10000):
    # sample a mini batch
    sample_index = torch.tensor(np.random.choice(N, batch_size))

    # Forward pass: compute predicted y by passing x to the model.
    y_pred = model(x[sample_index])

    # Compute the least-square loss.
    loss = loss_fn(y_pred, y[sample_index])


    # use the optimizer object to zero all of the gradients for the variables it will update,
    # i.e. the weights of the model. Checkout docs of torch.autograd.backward for more details.
    optimizer.zero_grad()

    # compute gradient of the loss with respect to model parameters (backward autodiff)
    loss.backward()

    # call the step function of the optimizer to make one update of the parameters
    optimizer.step()

y_pred = model(x)
plt.figure()
plt.plot(x.numpy(), y.numpy(), 'x')
plt.plot(x.numpy(), y_pred.detach().numpy(), 'rx')
```

```
# D_in is input dimension;
# H is dimension of the hidden layers; D_out is output dimension.
batch_size = 32
D_in, H, D_out = 1, 64, 1
```




Back to Q-learning

Q-learning with a table cannot work for high-dimensional spaces nor for continuous state/action spaces!

Idea: replace the table with a function approximator (e.g. a neural network) - still assume discrete number of actions

$$Q(x, u) \simeq Q(x, u, w)$$



Back to Q-learning

The problem can be written as a least square problem

We can compute the right side of Bellman equation
from data collected during one episode

$$y_t = g(x_t, u_t) + \alpha \min_a Q(x_{t+1}, a, w)$$

and then do one step of gradient descent on the weights
of the neural network to minimize the TD error

$$\min_w ||y_t - Q(x_t, u_t, w)||^2$$

Q-learning with a neural network

Initialize $Q(x, u, w)$ with random weights w

For each episode:

Choose an initial state x_0

Loop for each step of the episode:

Choose an action u_t using an ϵ -greedy policy from Q

Observe the next state x_{t+1}

Compute $y_t = g(x_t, u_t) + \alpha \min_a Q(x_{t+1}, a, w)$

Update the weights of the neural network by doing one iteration of stochastic gradient descent

$$\min_w ||y_t - Q(x_t, u_t, w)||^2$$

Back to Q-learning

Problem: a direct (naive) approach using solely current episode data tend to be unstable (i.e. it diverges):

- The sequence of observations are correlated
- Small changes in Q can lead to large changes in policy

Back to Q-learning



[Mnih et al., Nature, 2015]

Deep Q-network (DQN)

[Mnih et al., Nature, 2015]

Problem: a direct (naive) approach using solely current episode data tend to be unstable (i.e. it diverges):



- The sequence of observations are correlated
- Small changes in Q can lead to large changes in policy

Solution 1)

Use a “replay” memory of a previous samples from which we randomly sample the next training batch (remove correlations)

Solution 2)

Use 2 Q-networks to avoid correlations due to updates



Deep Q-network (DQN)

[Mnih et al., Nature, 2015]

Initialize replay memory D of size N

Initialize Q-network with random weights θ

Initialize target \hat{Q} function with weights $\theta^- = \theta$

For each episode:

Start from an initial state x_0

Loop for each step t of the episode:

Choose a control action u_t using Q (e.g. ϵ -greedy policy)

Do u_t and observe the next state x_{t+1}

Compute $y_t = g(x_t, u_t) + \alpha \min_a \hat{Q}(x_{t+1}, a, \theta^-)$! here we use the target network

Store (x_t, u_t, y_t, x_{t+1}) in memory D

Sample minibatch K of transitions (x_k, u_k, y_k, x_{k+1}) from D


Gradient descent on θ to minimize $\sum_K ||Q(x_k, u_k, \theta) - y_k||^2$

Every C steps reset the target network by setting $\theta^- = \theta$

Deep Q-network (DQN)

[Mnih et al., Nature, 2015]

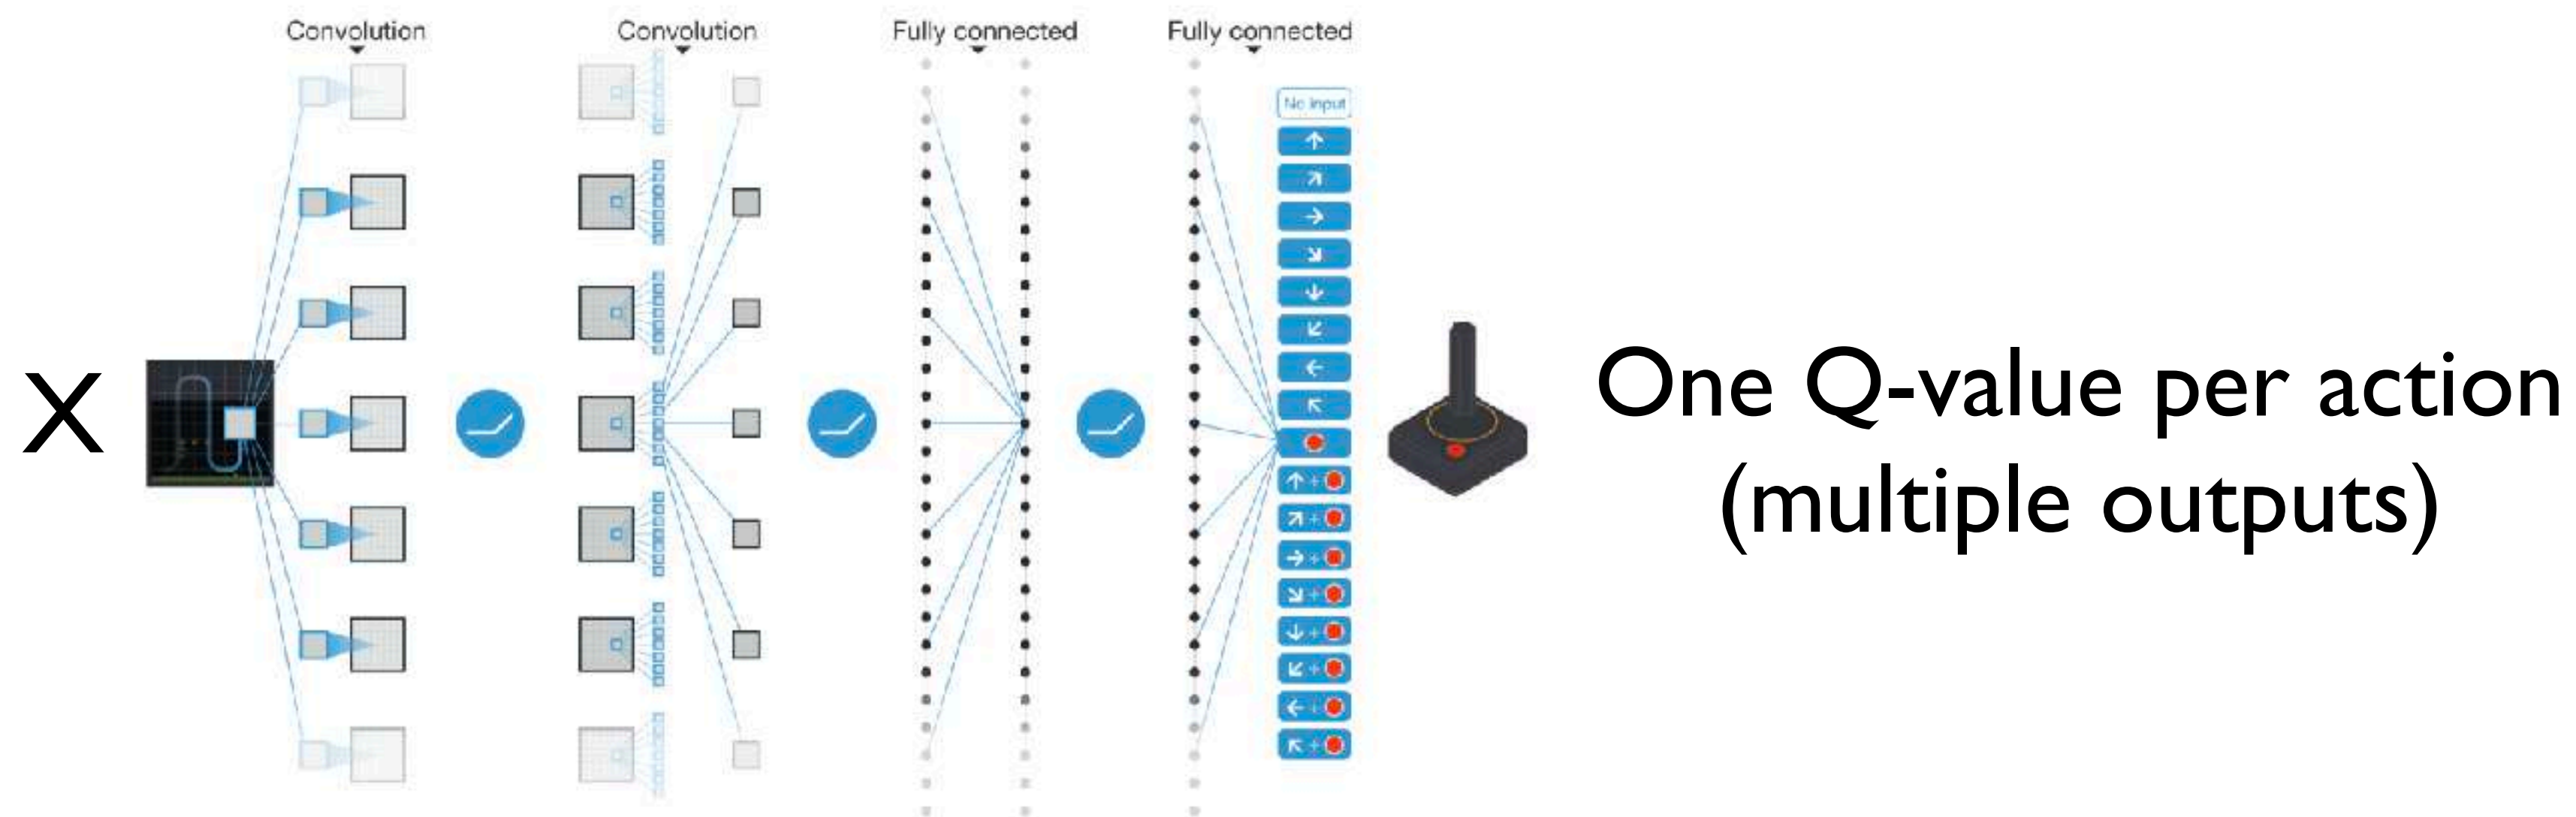

Pre-processing states of the system



Deep Q-network (DQN)

[Mnih et al., Nature, 2015]

Network Architecture



Deep Q-network (DQN)

[Mnih et al., Nature, 2015]

Training

49 games:

- a different Q network is used for each game
- same parameters for learning each game

mini-batches of size 32

ϵ -greedy with $\epsilon = 1$ at the beginning of learning and linearly decreases until $\epsilon = 0.1$ after first 1 million frames

trained on 50 million frames


=> 32 days of game experience in total!

(the human player was allowed only 2h of training)

replay memory size: 1 million samples (FIFO)


Deep Q-network (DQN)

[Mnih et al., Nature, 2015]




Deep Q-network (DQN)

[Mnih et al., Nature, 2015]



Now we can do Q-learning using continuous states and high dimensional inputs!

What about a continuous action space?



What about continuous action space?

Problem: we need to evaluate the min to be able to do Q-learning with a function approximator

$$||Q(x_t, u_t, \theta) - g(x_t, u_t) - \alpha \min_a \hat{Q}(x_{t+1}, a, \theta^-)||^2$$

Solution: use another neural network to approximate the min operator (i.e. to approximate the optimal policy)