

ROB-GY 6323
reinforcement learning and optimal
control for robotics

Lecture 13
Playing Go with Monte-Carlo Tree Search

Course material

All necessary material will be posted on Brightspace
Code will be posted on the Github site of the class

<https://github.com/righetti/optlearningcontrol>

Discussions/Forum with Slack

Contact

ludovic.righetti@nyu.edu

Office hours in person

Wednesday 3pm to 4pm

370 Jay street - room 801

Course Assistant

Armand Jordana

aj2988@nyu.edu

Office hours Monday 1pm to 2pm

Rogers Hall 515



any other time by appointment only

Schedule

Week	Lecture		Homework	Project	
1	<u>Intro</u>	Lecture 1: introduction			
2	<u>Trajectory optimization</u>	Lecture 2: Basics of optimization	HW 1		
3		Lecture 3: QPs			
4		Lecture 4: Nonlinear optimal control			
5		Lecture 5: Model-predictive control			
6		Lecture 6: Sampling-based optimal control	HW 2		
7	Lecture 7: Bellman's principle				
8	<u>Policy optimization</u>	Lecture 8: Value iteration / policy iteration		Project 1	
9		Lecture 9: Q-learning	HW 3		
10		Lecture 10: Deep Q learning			
11		Lecture 11: Actor-critic algorithms			
12		Lecture 12: Learning by demonstration	HW 4	Project 2	
13		Lecture 13: Monte-Carlo Tree Search			
14		Lecture 14: Beyond the class			
15	Finals week				

HW4 is due December 8th

Project 2 is due December 19th

Paper report (due December 19th - no deadline extension)

Goal: read one scientific paper and understand it

Pick one paper from the list of papers posted on brightspace

Report (maximum 2 pages - IEEE format double column)

It should contain 3 sections:

1. A section that summarizes the paper:
What was done? How was it done?
Why was it worth doing? What are the results?
2. A section explaining how the paper relates to the algorithms seen in class.
Which algorithms? What is different?
3. A section containing a critical discussion on the paper: pros and cons.
What seems to work and what convinces you about the result? What are the issues/limitations? What could be done better? What should be done next?

Do not copy equations or figures from the paper - keep your explanations to the point

Imitation learning

Learning optimal policies from demonstration

Behavioral cloning: learning policies from demonstrations

Provide a lot of demonstrations and learn a policy from it



Input: a dataset of demonstrations $(x_0, u_0, x_1, u_1, \dots, x_N, u_N)$

Output: a policy $u_n = \pi(x_n)$

A supervised learning problem!

$$\min_{\theta} \sum_N (\pi_{\theta}(x_n) - u_n)^2$$


Can learn very complex behaviors



[Chi et al. 2024]

Behavioral cloning: learning policies from demonstrations

Problem: compounding errors leads the robot out of demonstration distribution




Dataset aggregation DAGGER

Idea: as the robot does into “unseen territory” collect data and ask an “expert” to provide the correct control
(In effect we relabel the data the robot is collecting)

```
Initialize  $\mathcal{D} \leftarrow \emptyset$ .  
Initialize  $\hat{\pi}_1$  to any policy in  $\Pi$ .  
for  $i = 1$  to  $N$  do  
    Let  $\pi_i = \beta_i \pi^* + (1 - \beta_i) \hat{\pi}_i$ .  
    Sample  $T$ -step trajectories using  $\pi_i$ .  
    Get dataset  $\mathcal{D}_i = \{(s, \pi^*(s))\}$  of visited states by  $\pi_i$   
    and actions given by expert.  
    Aggregate datasets:  $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}_i$ .  
    Train classifier  $\hat{\pi}_{i+1}$  on  $\mathcal{D}$ .  
end for  
Return best  $\hat{\pi}_i$  on validation.
```

Algorithm 3.1: DAGGER Algorithm.



[Handa et al. 2022]

Learning cost functions from demonstrations

Inverse RL and apprenticeship learning



Inverse RL / inverse OC

Can we infer the cost function from a demonstration?

Useful for:

- Learning from demonstrations
- Apprenticeship learning
- Transferring skills across robots
- Also... analyzing human behavior

Playing the game of Go
A mixture of imitation learning, OC and RL



Deciding how to play with tree search

Go branching factor

For a game typically b^d number of moves to test
b is “breadth”, i.e. number of legal moves at each turn
d is the “depth”, i.e. the game length

For Go $b \sim 250$ and $d \sim 150$

Speeding up search: Monte-Carlo Tree Search

Invented by R. Coulom in 2006 (not new!)

4 steps to be repeated N times

1. Selection
2. Expansion
3. Simulation
4. Backup

Exploration vs. Exploitation (UCT)

AlphaGo 2016

Reduce the “breadth” and the “depth” of the search using MCTS

In addition, improve the sampling efficiency by:

- Learning a policy
- Learning a value function

Defining the states

Feature	# of planes	Description
Stone colour	3	Player stone / opponent stone / empty
Ones	1	A constant plane filled with 1
Turns since	8	How many turns since a move was played
Liberties	8	Number of liberties (empty adjacent points)
Capture size	8	How many opponent stones would be captured
Self-atari size	8	How many of own stones would be captured
Liberties after move	8	Number of liberties after this move is played
Ladder capture	1	Whether a move at this point is a successful ladder capture
Ladder escape	1	Whether a move at this point is a successful ladder escape
Sensibleness	1	Whether a move is legal and does not fill its own eyes
Zeros	1	A constant plane filled with 0
Player color	1	Whether current player is black

Step 1: Use supervised learning using human plays to learn a policy


Step 2: Use policy gradients to improve policy (using self-play)

Step 3: Use RL to compute value function of policy

Step 4: Monte-Carlo Tree Search using previously learned policy and value function to direct exploration

Stage I: learn a policy from Human players

Policy network



$$p_{\sigma}(a|s)$$


Policy learned with supervised learning
SL-policy

13 layers neural network -
accurate (57% / 55%) but slow
to evaluate (3ms)

$$p_{\pi}(a|s)$$

Policy with smaller network
- less accurate (24%) but fast
to evaluate (2us)

Stage 2: improve policy using RL policy gradient



p_{ρ} Policy learned using RL


$$\Delta\rho \propto \frac{\partial \log p_{\rho}(a_t | s_t)}{\partial \rho} z_t$$

$$z_t = \pm r(s_T)$$

RL policy won 80% games againsts SL policy

Stage 3: learning a value function

Value network



Self-play and randomization

$$v^p(s) = \mathbb{E}[z_t | s_t = s, a_{t \dots T} \sim p]$$


$$\Delta\theta \propto \frac{\partial v_{\theta}(s)}{\partial \theta} (z - v_{\theta}(s))$$

Stage 4: (modified) Monte-Carlo Tree Search

Each edge of the tree (action/state pair) stores an action value $Q(s,a)$, visit count $N(s,a)$ and prior probability $P(s,a)$

a

Selection




Go down the (partial) tree using

$$a_t = \operatorname{argmax}_a (Q(s_t, a) + u(s_t, a))$$

$$u(s, a) \propto \frac{P(s, a)}{1 + N(s, a)}$$

Stage 4: (modified) Monte-Carlo Tree Search


b Expansion



When a leaf is reached it can be expanded using the SL policy network

$$P(s, a) = p_\sigma(a|s)$$

Stage 4: (modified) Monte-Carlo Tree Search

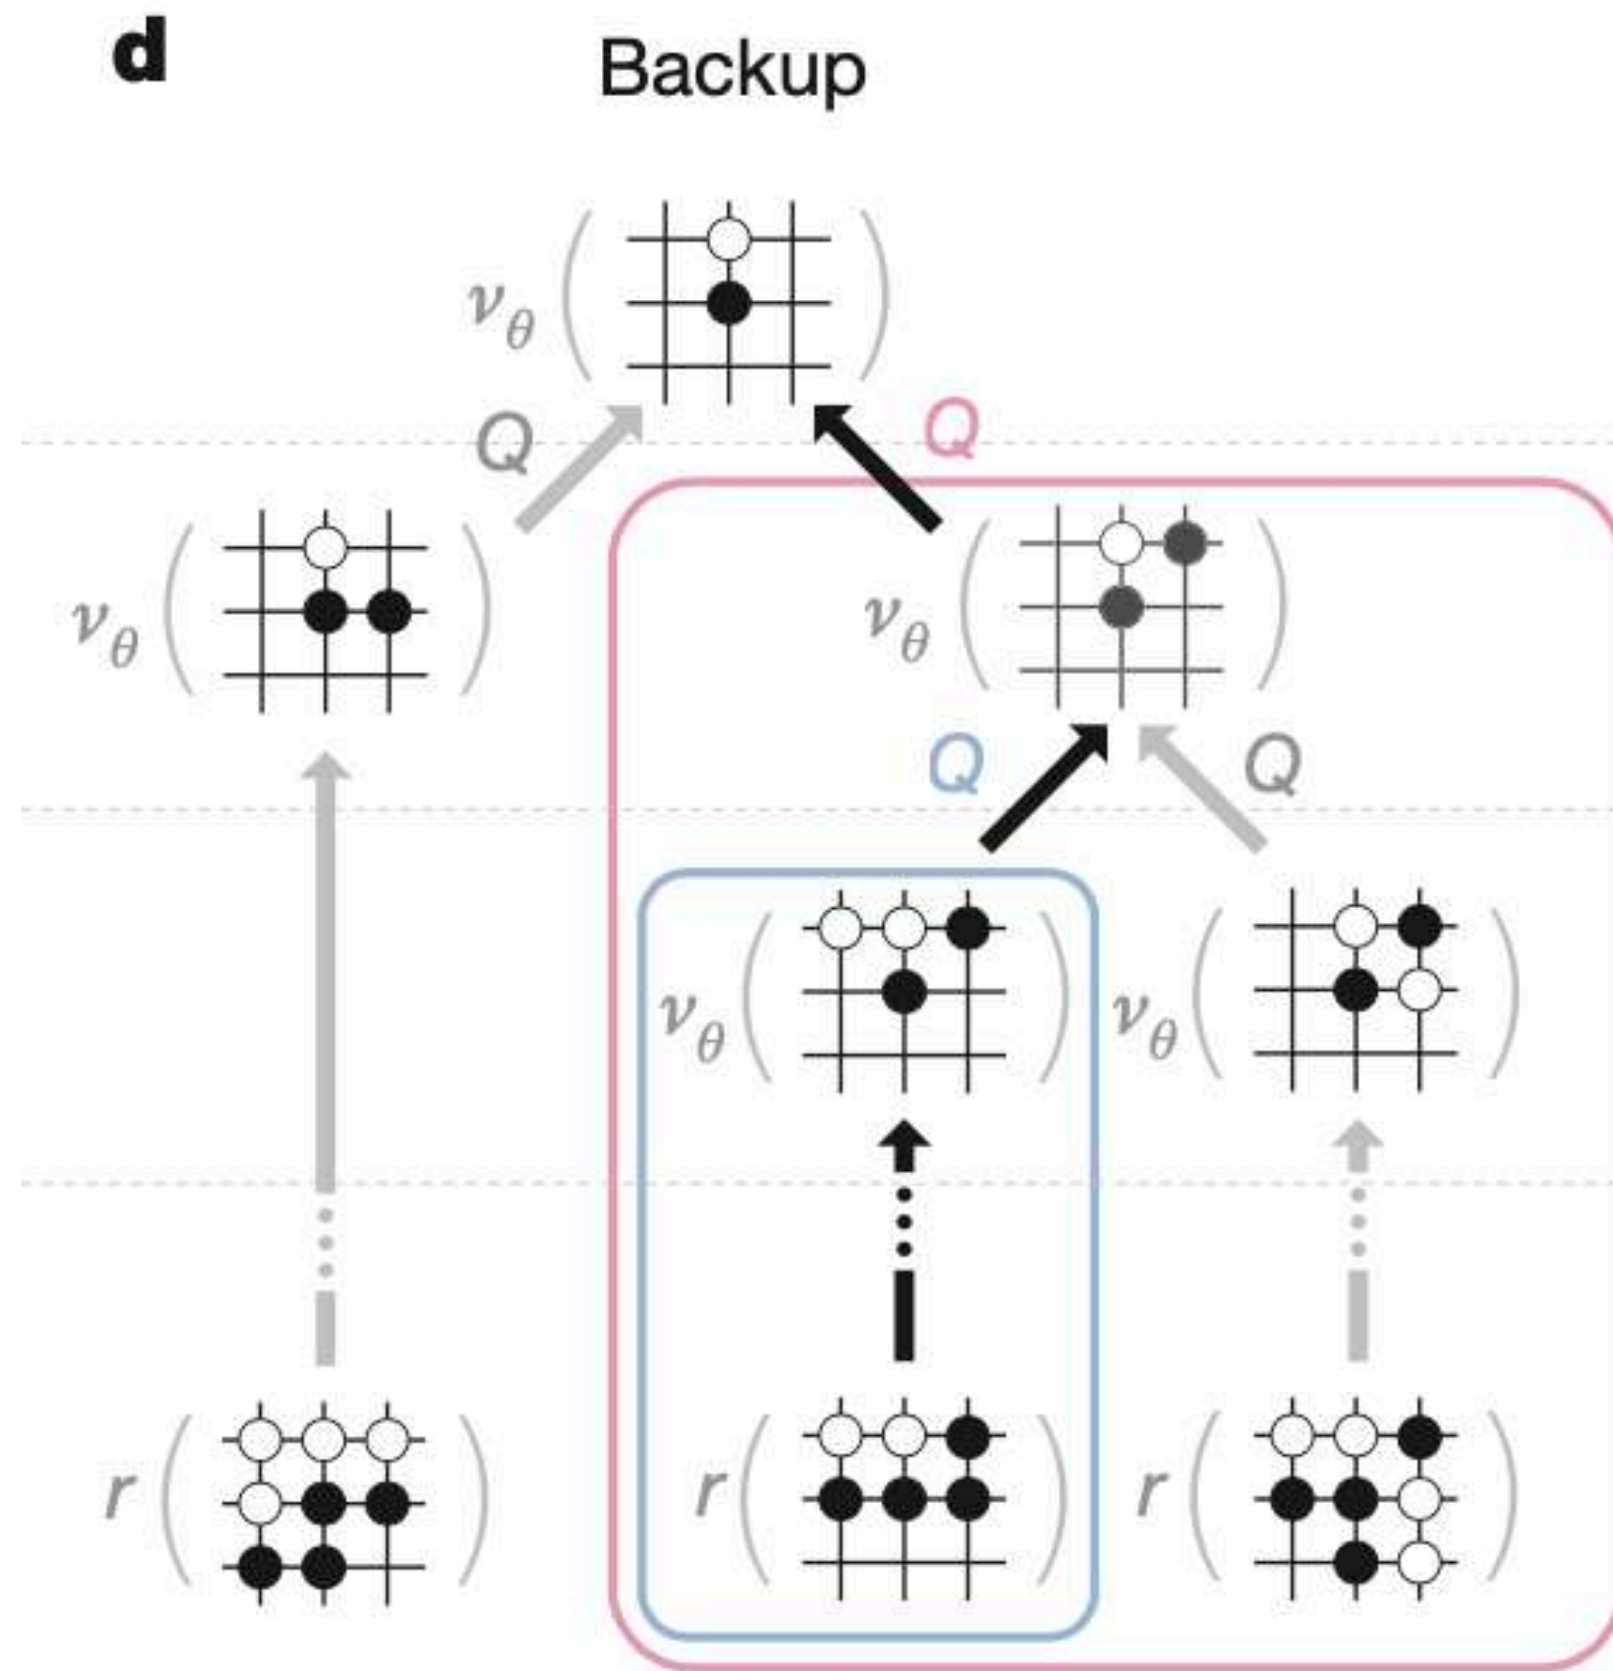


Evaluate the leaf node $V(s)$ using:

1. The learned value function $v^P(s)$
2. The outcome of a random simulated “play”

$$V(s_L) = (1 - \lambda)v_\theta(s_L) + \lambda z_L$$

Stage 4: (modified) Monte-Carlo Tree Search



$$N(s, a) = \sum_{i=1}^n 1(s, a, i)$$

$$Q(s, a) = \frac{1}{N(s, a)} \sum_{i=1}^n 1(s, a, i) V(s_L^i)$$

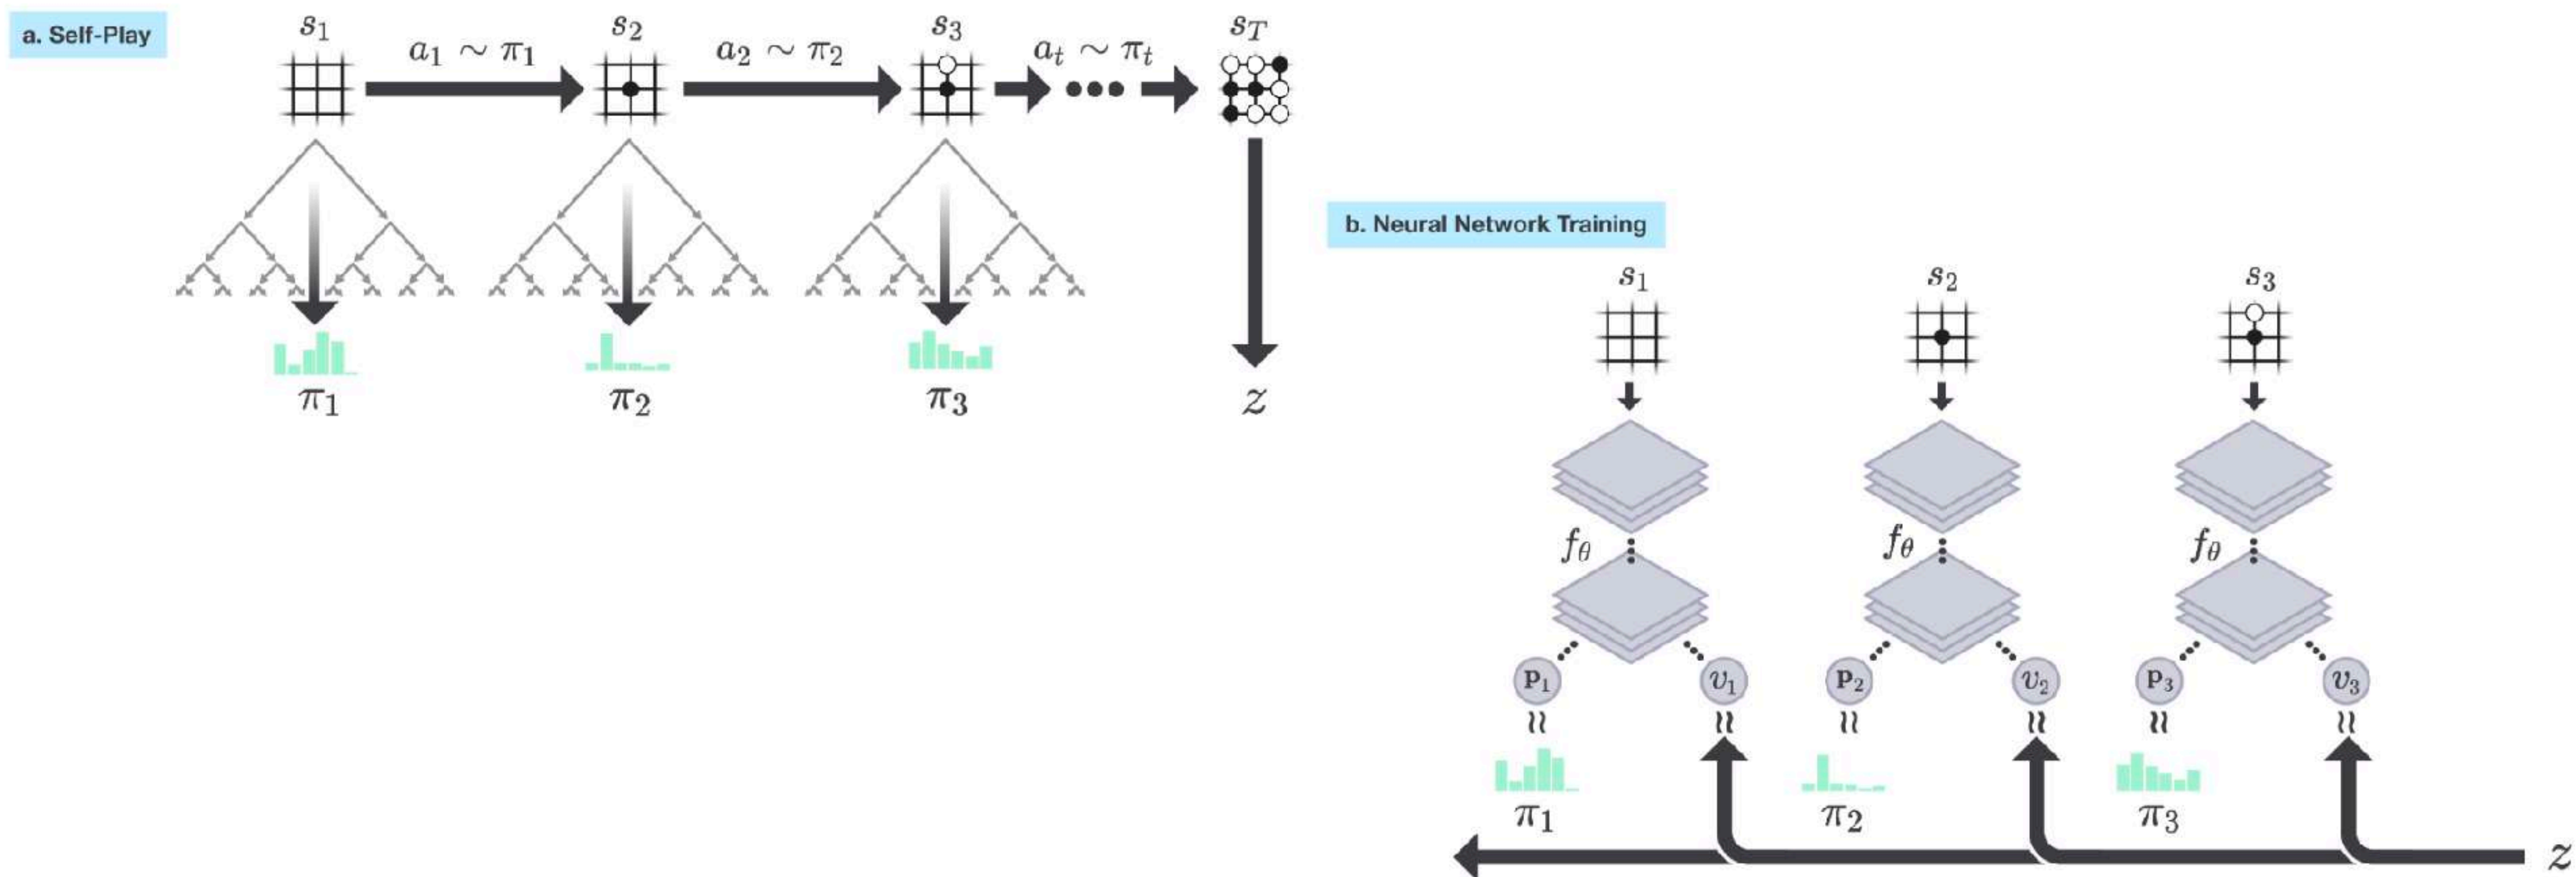
AlphaGoZero (2017)

Use self-play to learn a policy and value function (no SL)

Input features are only black and white stones (no other features)

Single NN for both policy and value

Simpler tree search - no evaluation through simulated play




AlphaGoZero (2017)

Use self-play to learn a policy and value function (no SL)

Input features are only black and white stones (no other features)

Single NN for both policy and value

Simpler tree search - no evaluation through simulated play



AlphaZero (2017)

Similar to AlphaGoZero but to play also Chess and Shogi

MuZero (2019)

Similar to AlphaGoZero but also learns the game model